OS 2018 v. 02

MIT ;)

https://pdos.csail.mit.edu/6.828/2018

0. POCITAC x86

Schema

- Von Neumann schema PC

- CPU
- Memory
- I/0
- zbernica

Program

- Pamat
  - Instrukcie
  - Udaje
- CPU
  - Interpretacia
  - Manipulacia

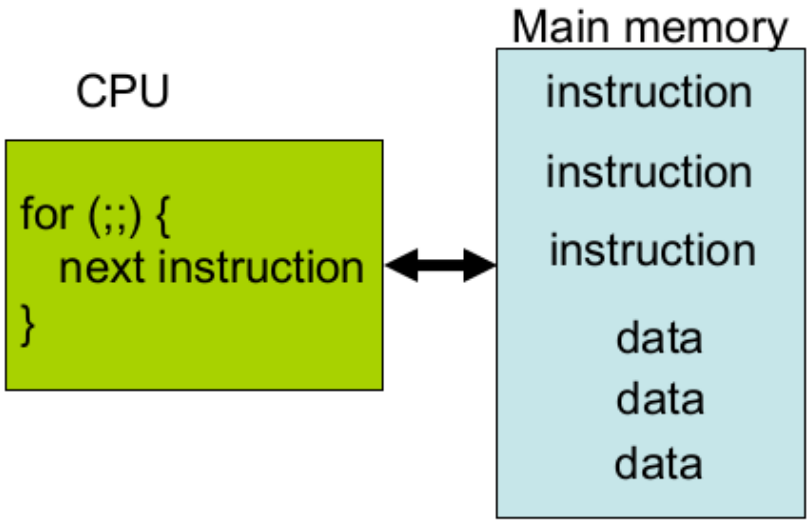

EIP

- Zvyseny po kazdej instrukcii
- Rozlicna dlzka instrukcii
- Automaticky modifikovany
  - Call
  - Ret
  - Jmp
- Neda sa priamo menit

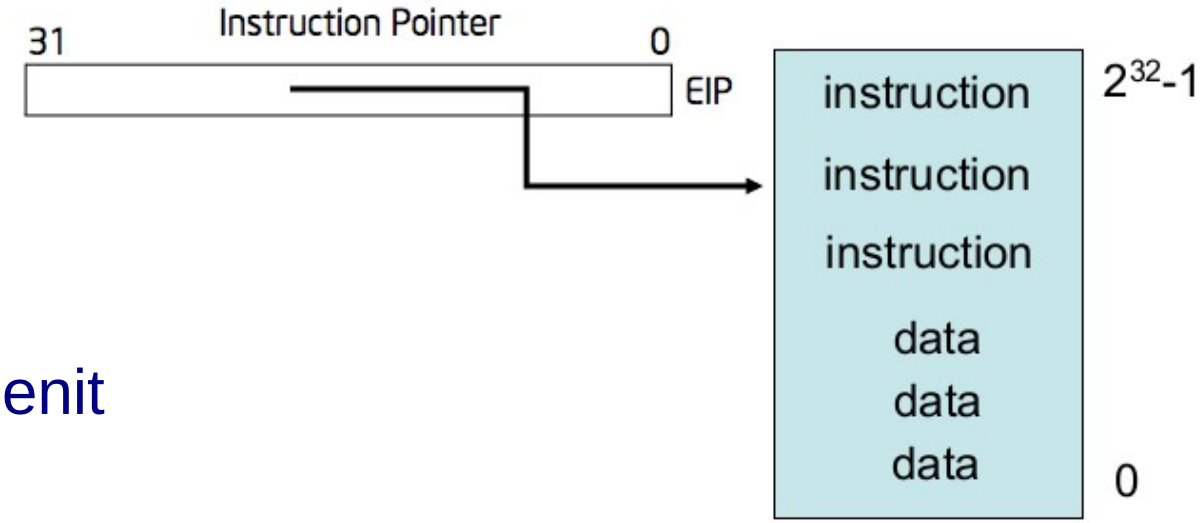

Registre CPU

- 8, 16, 32 bit
- Specialny ucel pre niektore
- Spatna kompatibilita

General-Purpose Registers

31 16	15 8	7 0	16-bit	32-bit
	AH	AL	AX	EAX
	BH	BL	BX	EBX
	CH	CL	CX	ECX
	DH	DL	DX	EDX
	BP			EBP
	SI			ESI
	DI			EDI
	SP			ESP

EFLAGS

- TEST
- CMP
- Conditional JMP

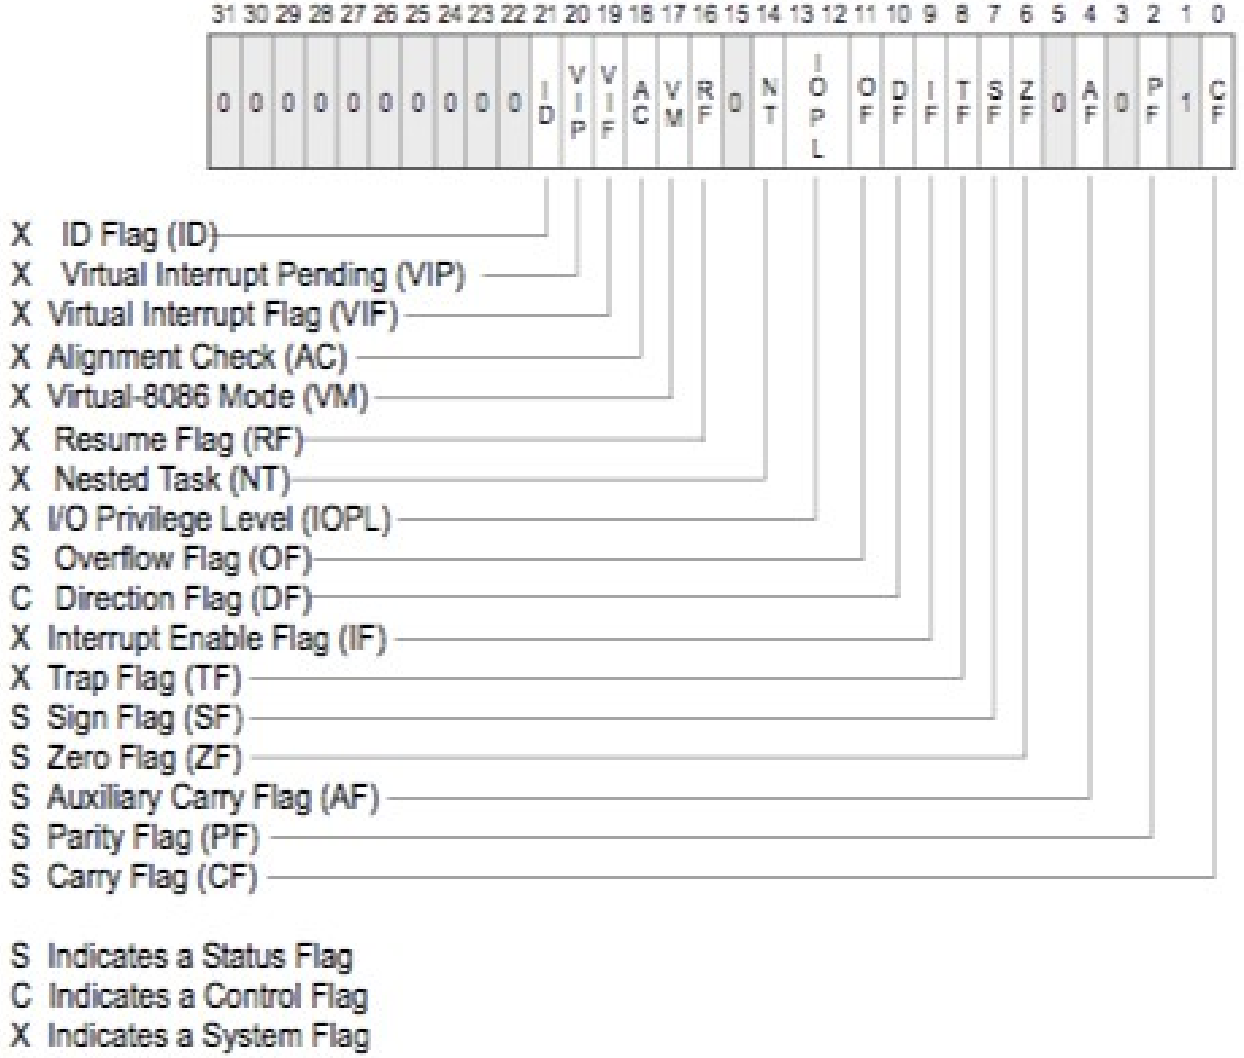

Zasobnikove instrukcie

- Push
- Pop
- Call
- Ret

Fyzicka pamat

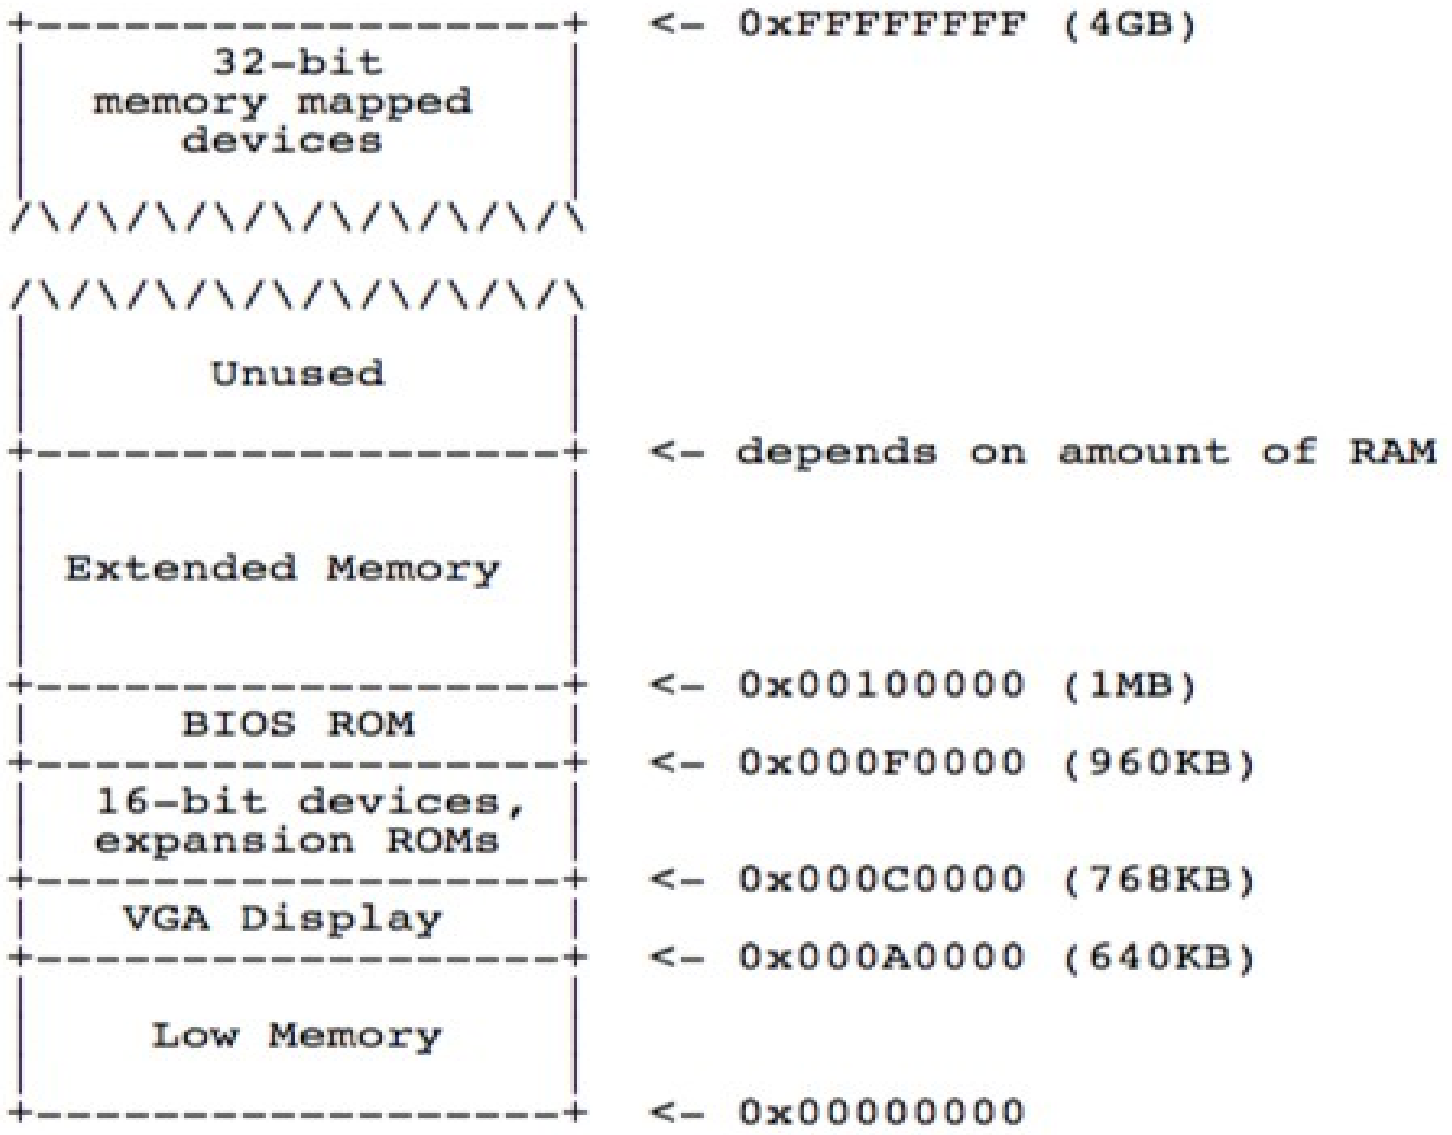

Instrukcna sada x86

- Prenos udajov (mov, push, pop, ...)
- Aritmetika (test, shl, add, sub, …)
- I/O (in, out, ...)
- Riadiace (jmp, jz, jnz, call, ret, …)
- Retazcove (rep, movsb, scasw, …)
- Systemove (iret, int, ...)

1. cast: ZAVEDENIE JADRA

BIOS

- Pripravit hw
- Najst bootovatelne medium
- Nacitat zavadzac z media na adresu 0x7c00
- Dat mu riadenia

Boot loader

- Prepnutie CPU do modernejsieho modu adresovania
- Nacitanie jadra do pamati
- Dat riadenie jadru

Boot loader - asm

- Cli – prerusenia hw, Cld – praca s retazcami...
- Realny mod CPU
  - 20 bitove adresy z 2 16 bitovych reg
  - CS, DS, ES, SS
  - Logicka, linearna a fyzicka adresa
  - Virtualne adresy

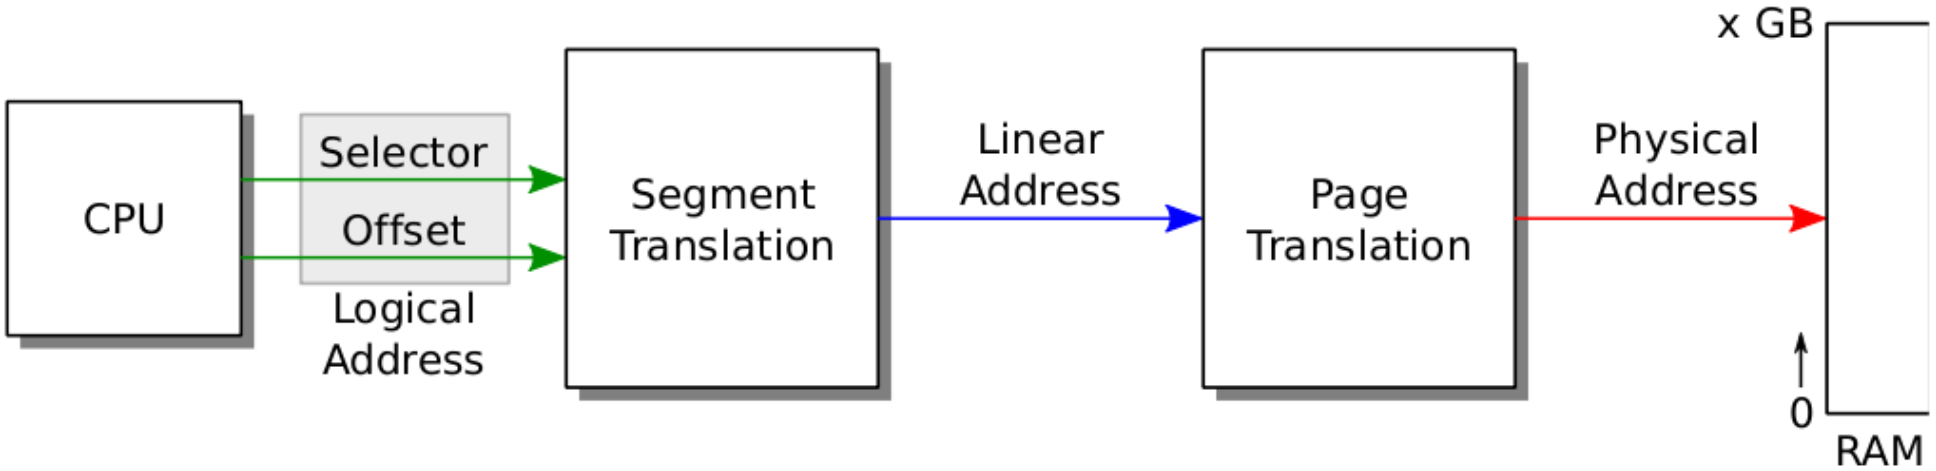

Boot loader - asm

- Obmedzenie realneho modu adresacie
  - 64kB per program
  - 1MB per fyzicka pamat
- Trik A20 pre zapnutie ostatnych adresnych bitov CPU!
  - BIOS
  - Klavesnicovy radic (JOS)
  - Tzv. “rychle” zapnutie cez port

Boot loader - asm

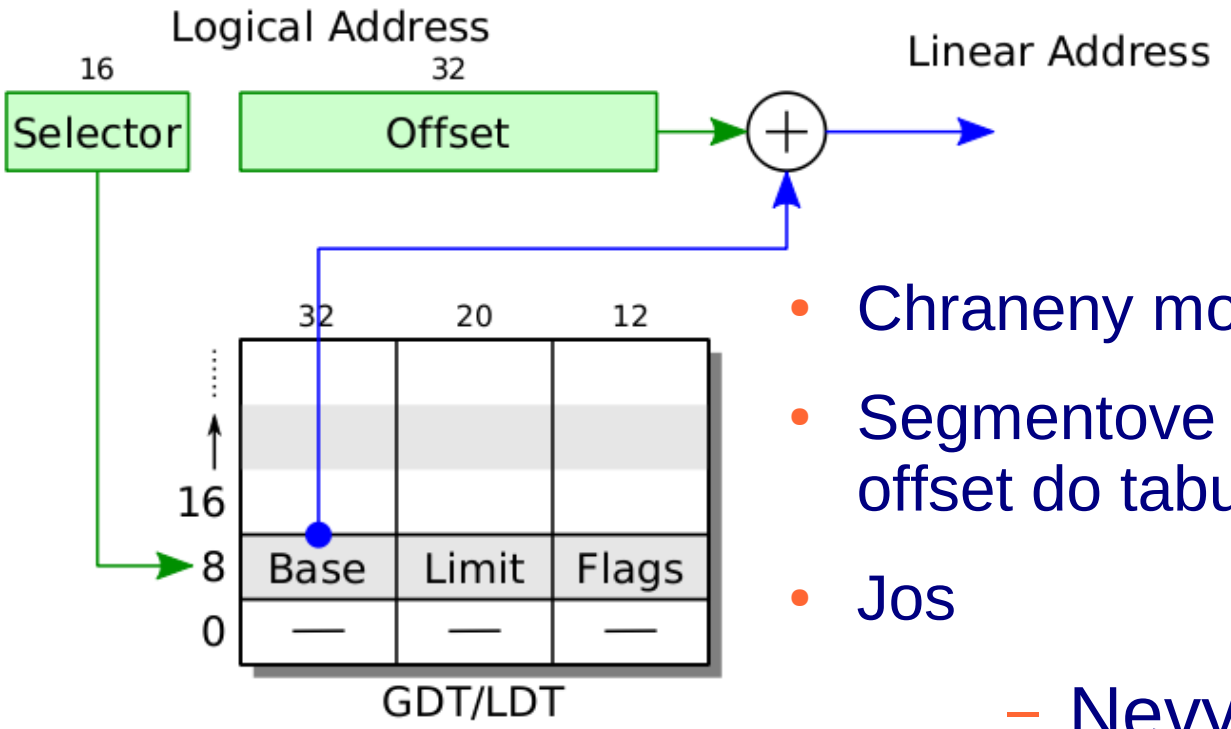

- Chraneny mod CPU
- Segmentove registre (selector) – offset do tabulky (base, limit, flags)
- Jos
  - Nevyuziva preklad cez tabulku
  - 3 polozky ma tabulka

Boot loader – asm a C

- Nastavenie zasobnika
- Spustenie bootmain()
  - Nacitat a spustit jadro
  - Co v pripade, ze 1. sektor neobsahuje kod?
  - Jadro: elf format, od 2. sektora na disku
  - Nacitanie sekcii do pamate
  - Spustenie jadra

Jadro

- Jadro je v pamati od 0x100000, nie od 0xf0100000
- Musi sa zapnut preklad adries (strankovanie)
- Mapuju sa dva regiony!!! Preco?

- Nastavenie registra ebp pre spravny “stack trace”
- Nastavenie zasobnika (registra esp)
- Spustenie kodu v C - funkcie i386_init()

2. cast: ZASOBNIK na x86

Konvencia volani - cdecl

```
int volany(int, int, int);

int volajuci(void) {
    int ret;
    ret = volany(1, 2, 3);
    ret += 5;
    return ret;
}
```

```
volajuci:
    ; novy stack frame
    push ebp
    mov ebp, esp
    ; uloz argumenty
    push 3
    push 2
    push 1
    ; zavolaj 'volaneho'
    call volany
    ; zrus argumenty
    add esp, 12
    ; pripocitaj 5
    add eax, 5
    ; obnov povodny stack frame
    pop ebp
    ; navrat
    ret
```

Zasobnik

```
int foobar(int a, int b, int c)
{
    int xx = a + 2;
    int yy = b + 3;
    int zz = c + 4;
    int sum = xx + yy + zz;

    return xx * yy * zz + sum;
}

int main()
{
    return foobar(77, 88, 99);
}
```

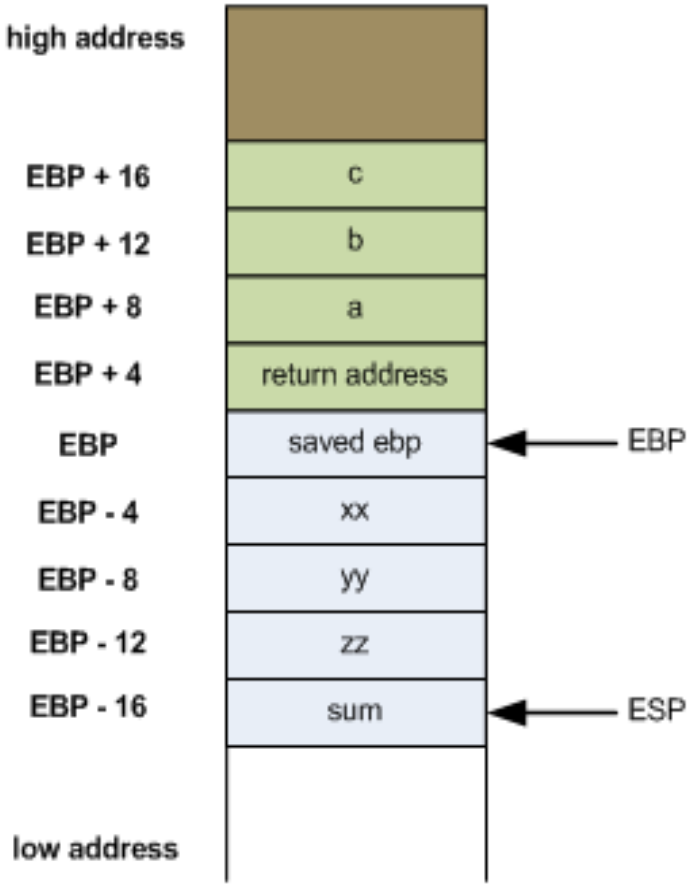

Zasobnik a volanie funkcie

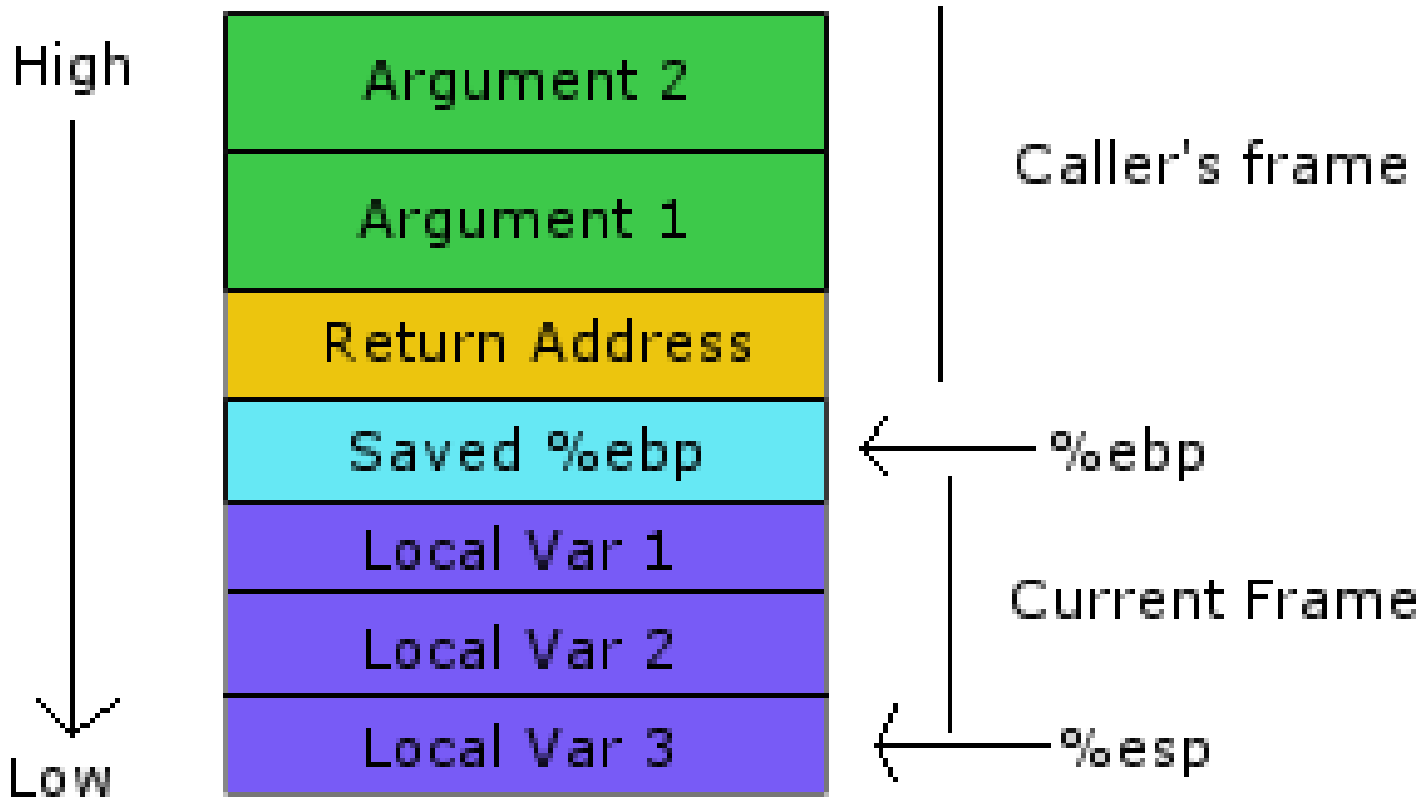

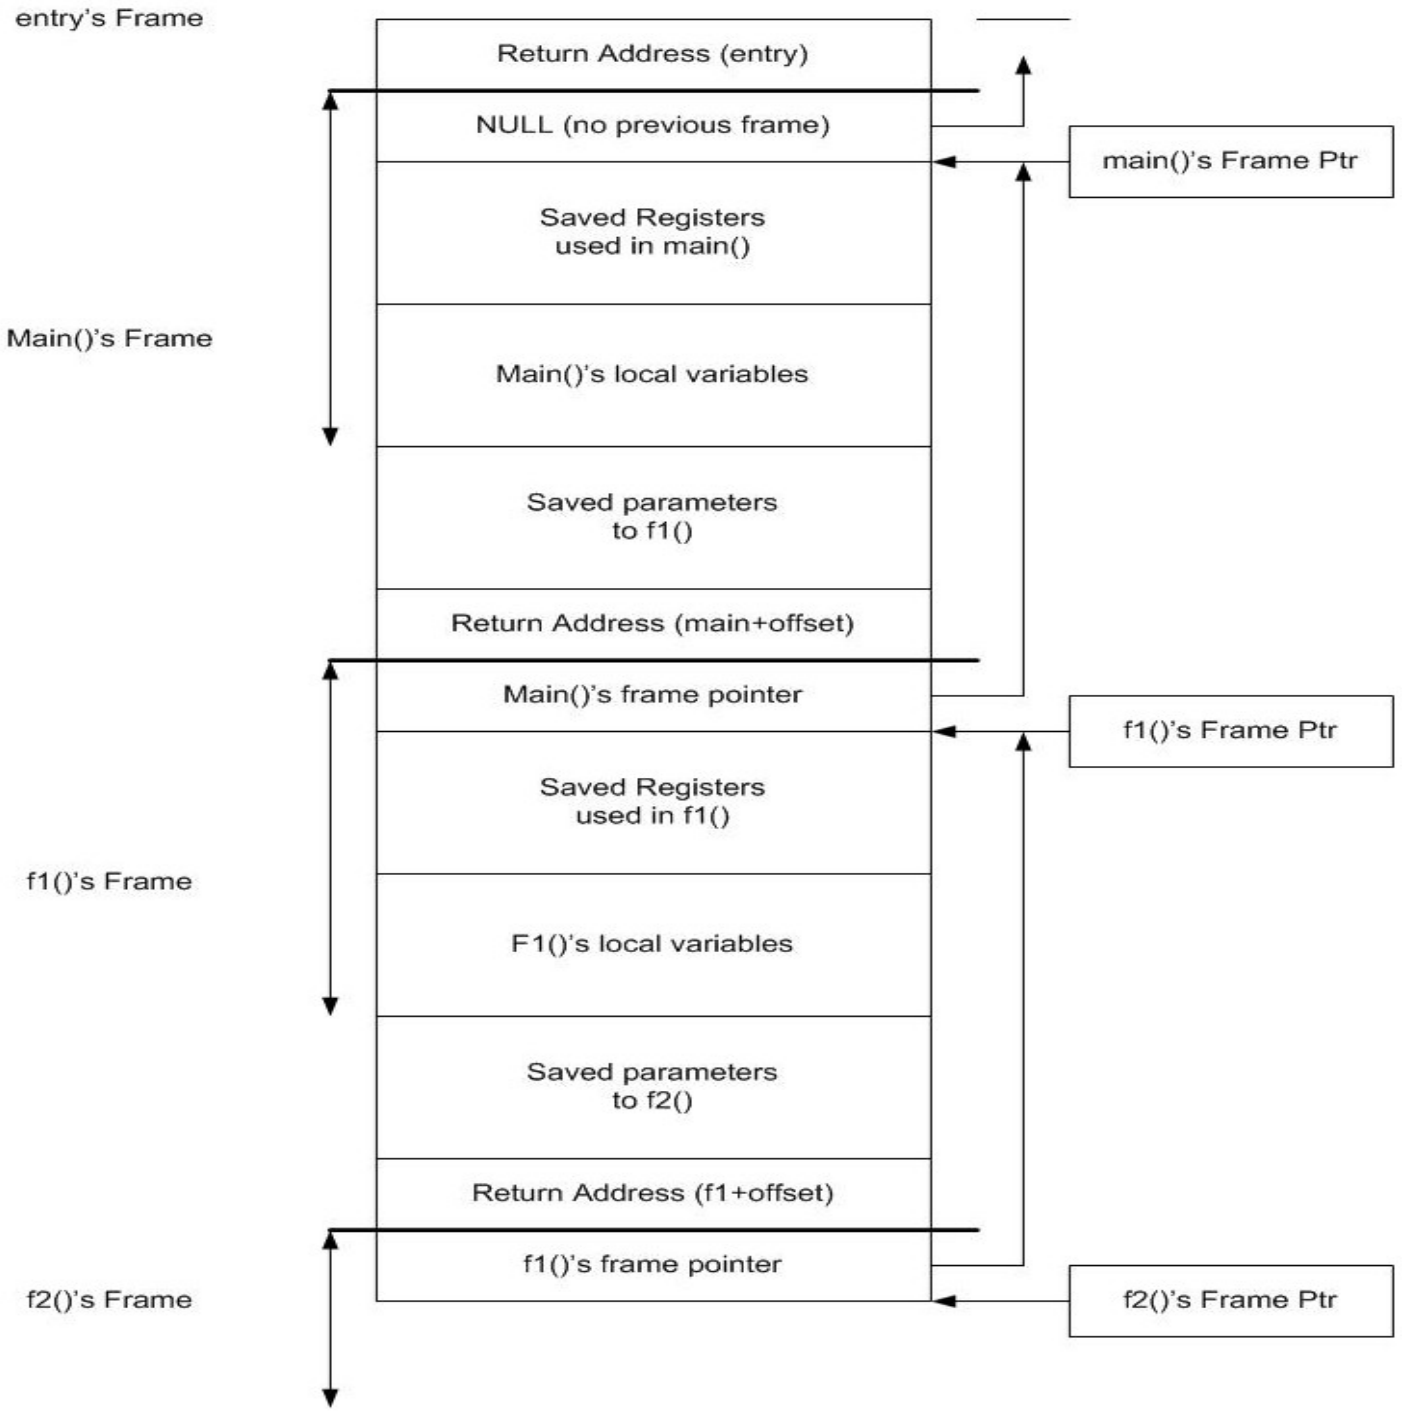
entry's Frame
Return Address (entry)
NULL (no previous frame)
main()'s Frame Ptr
Saved Registers
used in main()
Main()'s Frame
Main()'s local variables
Saved parameters
to f1()
Return Address (main+offset)
Main()'s frame pointer
f1()'s Frame Ptr
Saved Registers
used in f1()
f1()'s Frame
F1()'s local variables
Saved parameters
to f2()
Return Address (f1+offset)
f1()'s frame pointer
f2()'s Frame
f2()'s Frame Ptr

Prolog a Epilog funkcie

```
int volany(int, int, int);

int volajuci(void) {
    int ret;
    ret = volany(1, 2, 3);
    ret += 5;
    return ret;
}
```

```
volajuci:
    ; novy stack frame
    push ebp
    mov ebp, esp
    ; uloz argumenty
    push 3
    push 2
    push 1
    ; zavolaj 'volaneho'
    call volany
    ; zrus argumenty
    add esp, 12
    ; pripocitaj 5
    add eax, 5
    ; obnov povodny stack frame
    pop ebp
    ; navrat
    ret
```

Dalsi obrazok

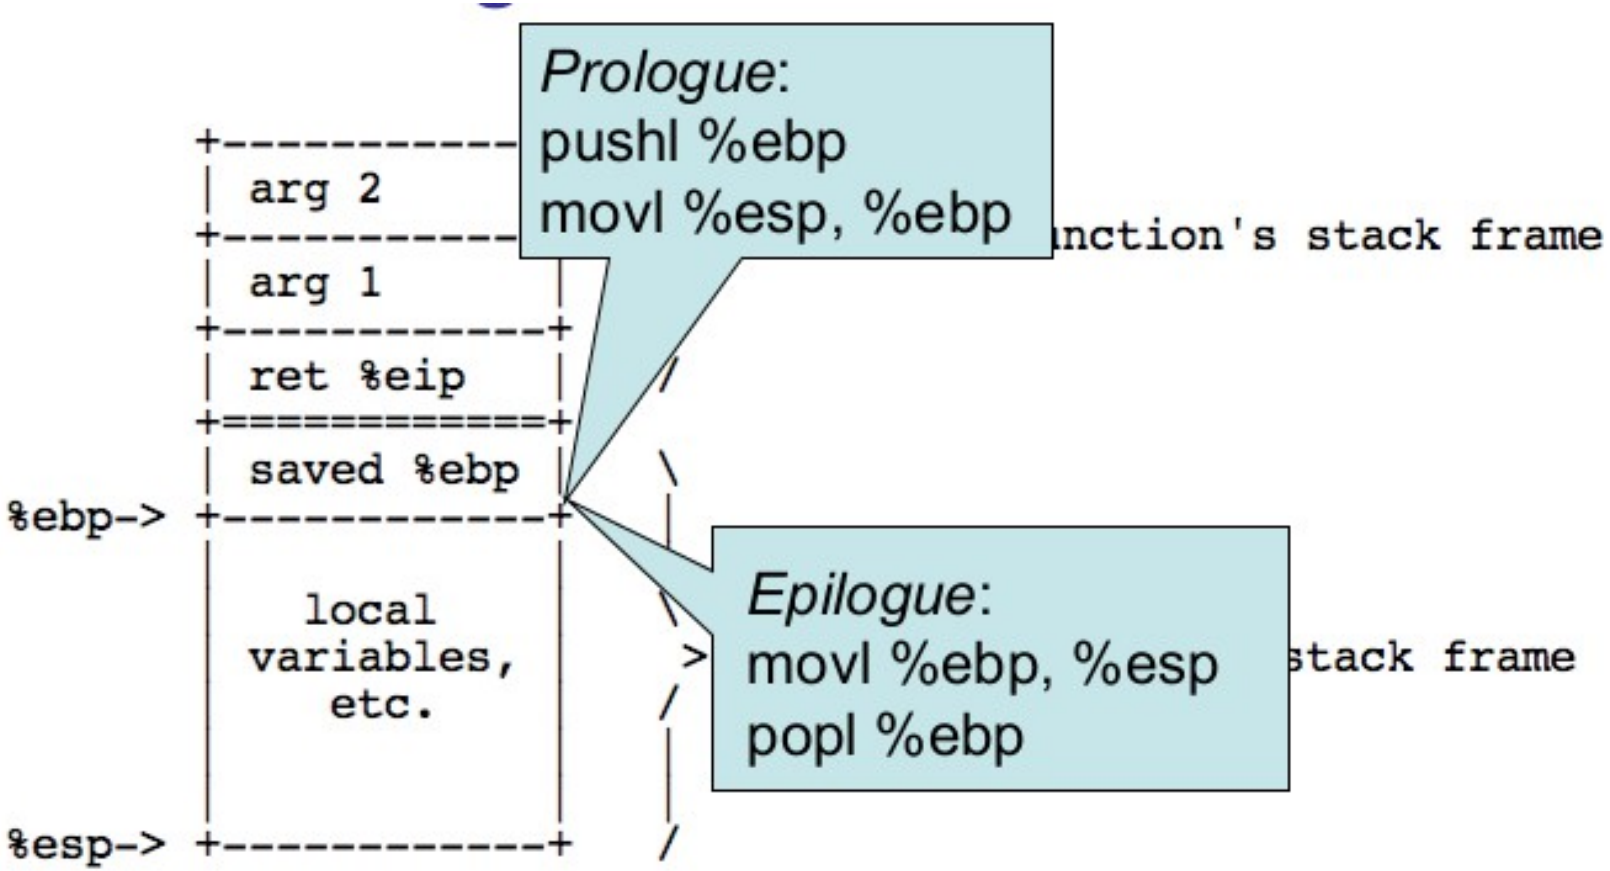

Priklad

```
int main(void) { return f(8)+1; }
int f(int x) { return g(x); }
int g(int x) { return x+3; }
```

```
_main:
                        prologue
        pushl %ebp
        movl %esp, %ebp
                        body
        pushl $8
        call _f
        addl $1, %eax
                        epilogue
        movl %ebp, %esp
        popl %ebp
        ret
_f:
                        prologue
        pushl %ebp
        movl %esp, %ebp
                        body
        pushl 8(%esp)
        call _g
                        epilogue
        movl %ebp, %esp
        popl %ebp
        ret

_g:
                        prologue
        pushl %ebp
        movl %esp, %ebp
                        save %ebx
        pushl %ebx
                        body
        movl 8(%ebp), %ebx
        addl $3, %ebx
        movl %ebx, %eax
                        restore %ebx
        popl %ebx
                        epilogue
        movl %ebp, %esp
        popl %ebp
        ret
```

test_backtrace() v JOS

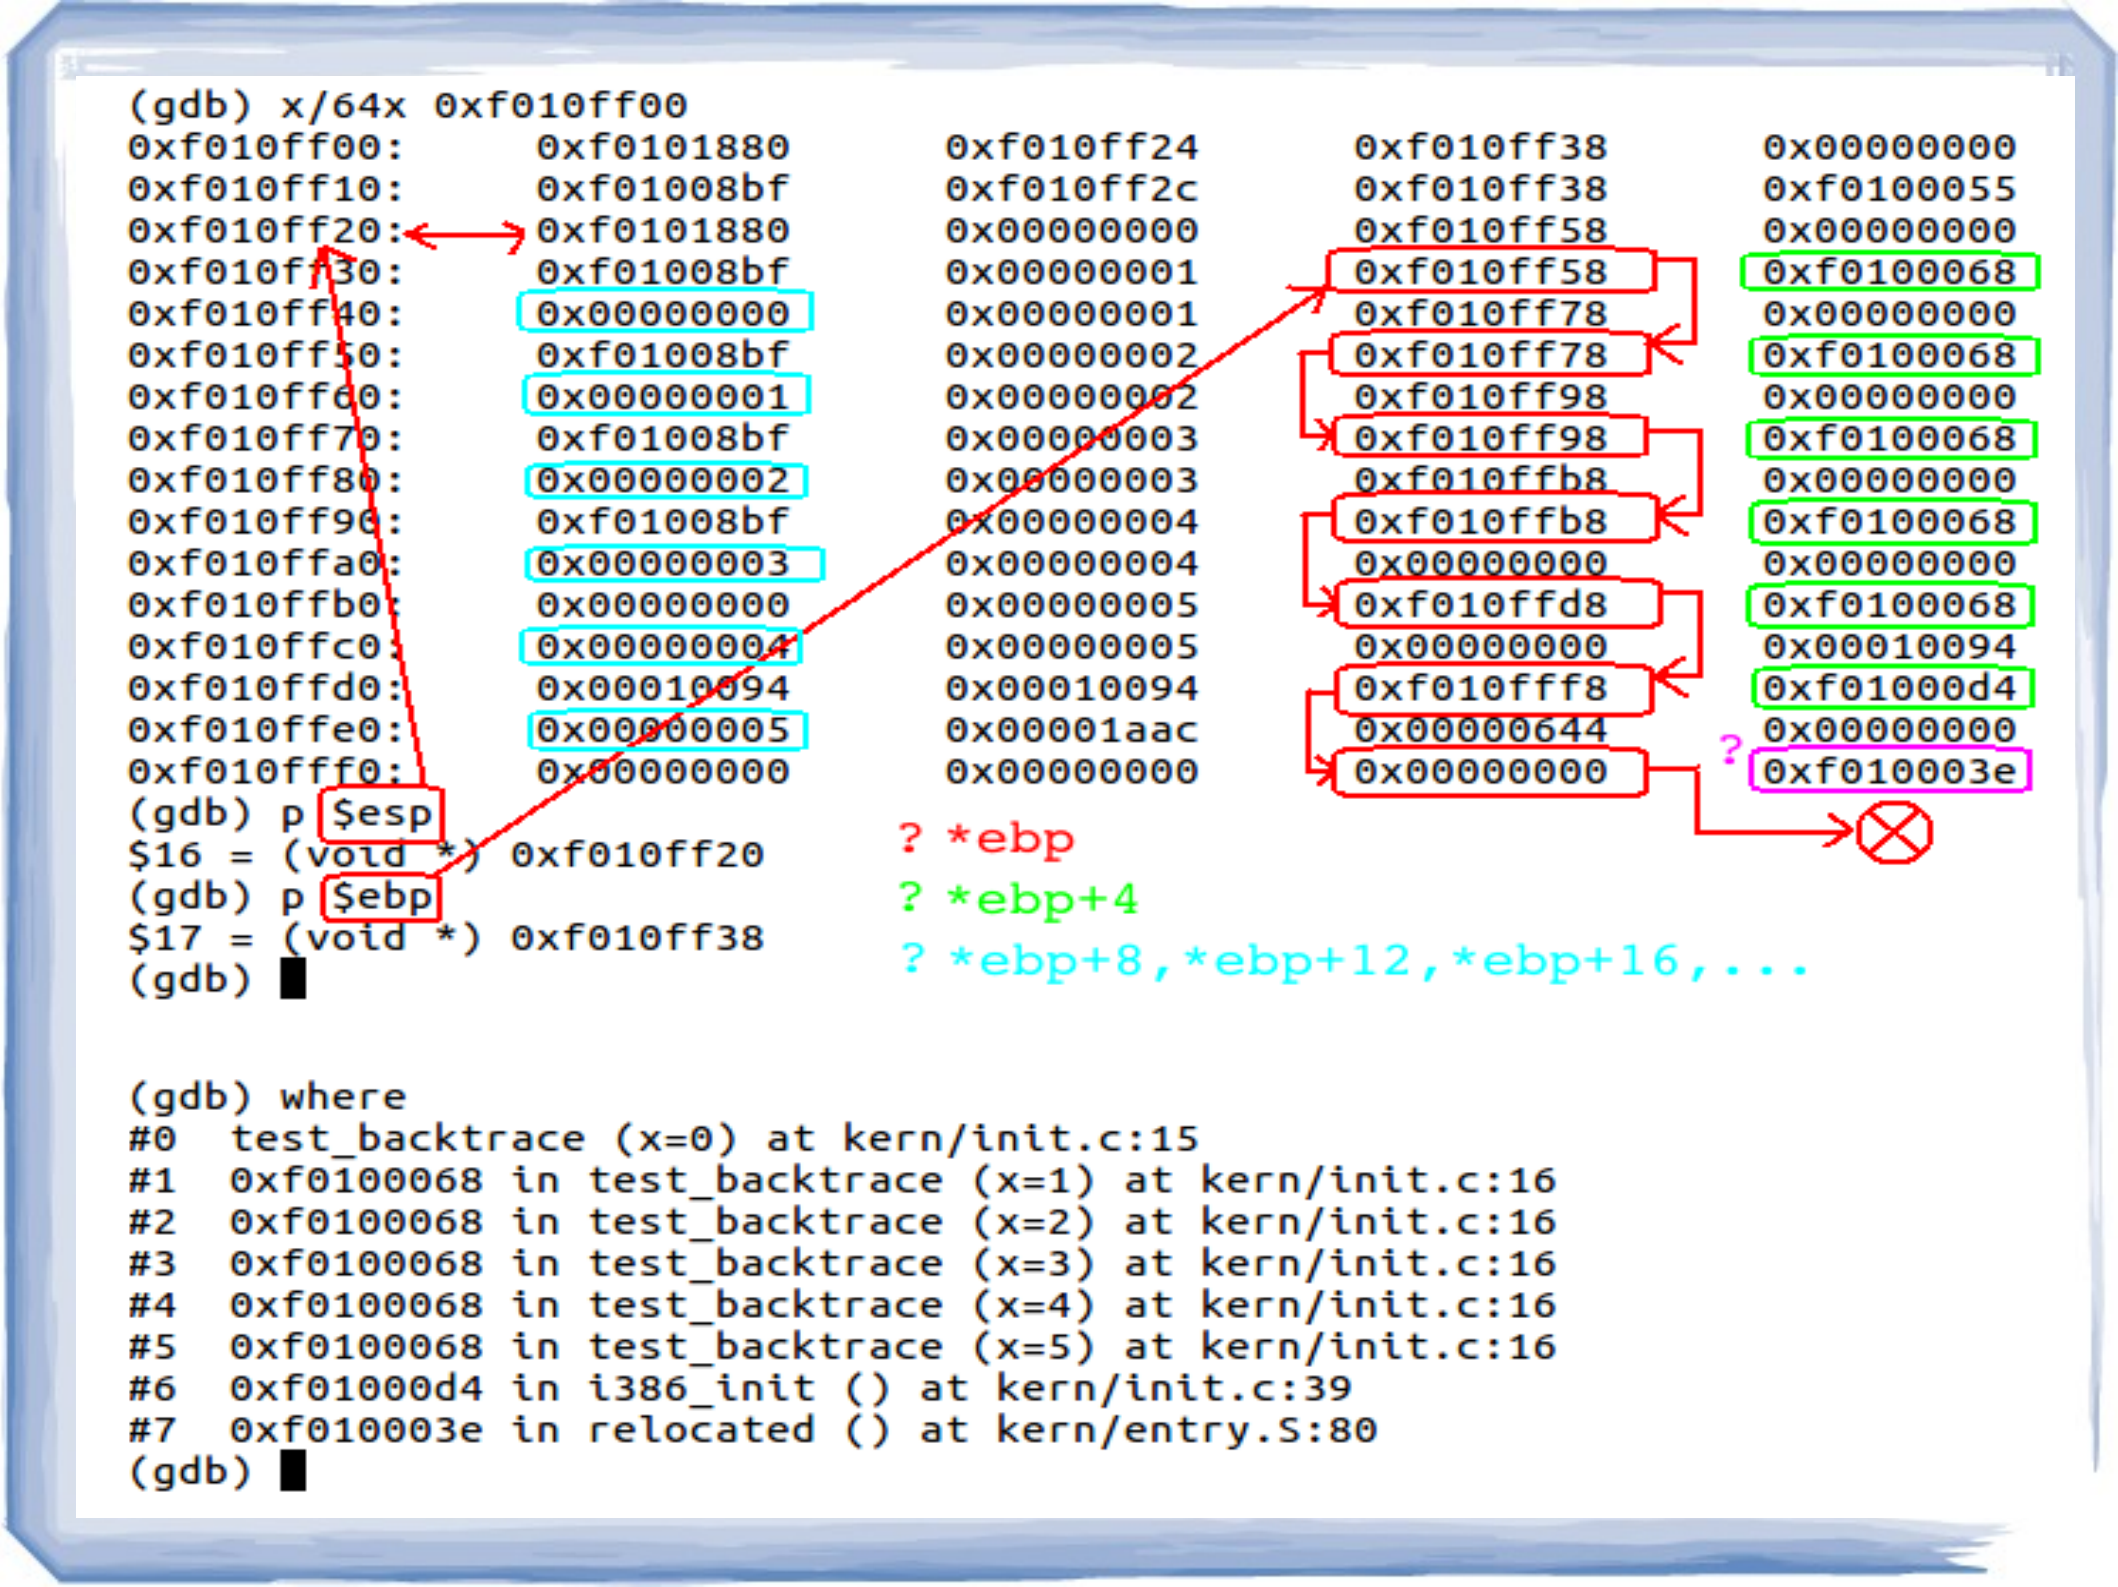


```
(gdb) where
#0  test_backtrace (x=0) at kern/init.c:15
#1  0xf0100068 in test_backtrace (x=1) at kern/init.c:16
#2  0xf0100068 in test_backtrace (x=2) at kern/init.c:16
#3  0xf0100068 in test_backtrace (x=3) at kern/init.c:16
#4  0xf0100068 in test_backtrace (x=4) at kern/init.c:16
#5  0xf0100068 in test_backtrace (x=5) at kern/init.c:16
#6  0xf01000d4 in i386_init () at kern/init.c:39
#7  0xf010003e in relocated () at kern/entry.S:80
(gdb)
```